

Enhanced Lung Nodule Detection Through Integration of U-Net++ and Multi-Scale Gradual Integration CNN

Zhaoxu Zheng^{1*}, Kexin Zhang²

¹Department of Computer Science, Harbin University of Science and Technology, Harbin, China

²School of Cyber Science and Engineering, Nanjing University of Science and Technology, China

*0zcharles318@gmail.com

Abstract. Early detection of pulmonary nodules is crucial for improving lung cancer survival rates. While conventional deep learning approaches employing U-Net and 3D-CNN have shown promise, they often suffer from high false-positive rates and limited segmentation accuracy. This paper presents an enhanced two-stage detection pipeline that integrates U-Net++ for precise candidate generation with Multi-scale Gradual Integration CNN (MGI-CNN) for robust false-positive reduction. The proposed architecture leverages U-Net++'s nested skip pathways and deep supervision to capture fine-grained nodule boundaries, while MGI-CNN's multi-scale 3D feature integration effectively discriminates true nodules from anatomical structures. Experimental results on the LUNA16 dataset demonstrate that our pipeline achieves superior performance compared to traditional U-Net+MGI-CNN methods, with nodule-level recall improving from 93.4% to 97.8% and classification AUC increasing from 0.852 to 0.901. The synergistic combination of advanced segmentation and classification techniques represents a significant advancement in automated pulmonary nodule detection.

Keywords: Lung nodule detection, U-Net++, MGI-CNN, Deep learning, Medical image segmentation, False-positive reduction

1 Introduction

Lung cancer remains the leading cause of cancer-related mortality worldwide, with approximately 2.2 million new cases (11.4%) and 1.8 million deaths (18.0%) reported globally in 2020 [1]. Early detection through accurate identification of pulmonary nodules—small round or oval-shaped abnormal tissue growths in the lung parenchyma, typically measuring 3–30 mm in diameter—is critical for improving patient outcomes and for designing screening pathways in clinical practice [6, 7]. Computed tomography (CT) imaging serves as the primary modality for nodule detection due to its high spatial resolution and ability to capture detailed three-dimensional lung structures, but clinically useful systems must balance sensitivity with false-positive (FP) control to be deployable at scale [5, 14].

However, manual interpretation of CT scans faces significant challenges. The process is highly dependent on reader expertise and prone to inter-observer variability; small nodules, particularly those under 6 mm, are frequently missed when adjacent to vessels, pleura, or low-contrast regions. As LDCT screening expands, workload and variability motivate automated pipelines that maintain high recall without inflating FPs [5, 12, 13]. The challenge is compounded by subtle appearance differences between early malignant nodules and benign look-alikes (e.g., vessels in cross-section, lymph nodes, scars, inflammatory foci), which argues for architectures that encode multi-scale 3D context and attention to fine morphology [8, 10].

Deep learning has emerged as a transformative technology in medical image analysis, with convolutional neural networks (CNNs) demonstrating remarkable success in automated nodule detection. The seminal U-Net architecture [2], with its symmetric encoder-decoder structure and skip connections, established a new paradigm for medical image segmentation. Contemporary surveys further evaluate U-Net’s strengths and weaknesses in small-object, low-contrast regimes typical of lung nodules and recommend deeply supervised, modernized variants to improve boundary fidelity and optimization stability [12–14].




Fig. 1. Traditional U-Net architecture showing encoder-decoder structure with skip connections. The contracting path captures context while the expanding path enables precise localization.

As illustrated in Fig. 1, U-Net uses a contracting path to capture context and an expanding path for precise localization [2]. Despite its success, U-Net-based approaches still face limitations for lung nodule detection: the architecture often struggles with varying nodule sizes, irregular shapes, and subtle texture differences. More critically, these methods can generate numerous false positives by mistaking blood vessels, lymph nodes, and other anatomical structures for nodules [5, 14].

To address this challenge, researchers have explored two complementary families of detectors. One-stage 3D networks couple proposal generation and classification in a single pass [9], whereas two-stage cascades first seek high-recall candidates and then apply targeted FP reduction—an attractive strategy for screening because it

decouples sensitivity from specificity and enables explicit volumetric reasoning downstream [5, 16]. Architecturally, multi-scale and attention mechanisms have proved effective for thoracic CT [8], and cross-domain 3D studies show that multi-scale integration blocks can enhance volumetric discriminability—an idea directly applicable to nodule analysis [10]. In particular, the Multi-scale Gradual Integration CNN (MGI-CNN) fuses patches across physical scales via bidirectional (zoom-in/zoom-out) streams to couple fine morphology with broader anatomic context, achieving strong FP reduction on LUNA16 [4].

Recent advances in network design have also improved the candidate-generation stage. UNet++ introduces nested (dense) skip pathways and deep supervision to bridge the semantic gap between encoder and decoder, yielding more reliable boundaries on precision-critical biomedical targets, including lung nodules [3]. Further gains come from modern encoders and fine-tuning strategies—e.g., Nested U-Net with EfficientNet and carefully tuned training protocols improve mask fidelity on dense, small objects—supporting the use of stronger proposal backbones in cascaded pipelines [11,14,15]. Reviews in radiology echo these methodological trends, recommending the pairing of high-sensitivity, deeply supervised segmentation with context-aware 3D classification to balance sensitivity and specificity in LDCT workflows [5, 7].

This paper presents a lung nodule detection framework that follows this guidance by synergistically integrating UNet++ for high-recall, boundary-aware candidate generation with an MGI-CNN-based 3D classifier for robust FP reduction [3, 4]. Our design aligns with evidence favoring deeply supervised segmentation, multi-scale volumetric context, and screening-ready specificity control [8, 12, 13], while remaining practical for deployment in LDCT programs [5, 7]. Our key contributions are:

1. An enhanced two-stage detection pipeline that replaces traditional U-Net with U-Net++ for improved candidate generation, demonstrating measurable improvements in nodule-level recall (93.4% to 97.8%) and segmentation quality metrics.
2. Comprehensive experimental validation on the LUNA16 dataset showing that better segmentation quality from U-Net++ directly translates to improved false-positive reduction performance, with MGI-CNN achieving an AUC increase from 0.852 to 0.901.
3. Detailed analysis of training dynamics and convergence behavior, demonstrating that UNet++'s nested skip connections and deep supervision provide more stable optimization and better generalization compared to traditional U-Net.

2 Methods

2.1 Stage 1: Enhanced Candidate Generation with U-Net++

U-Net++ fundamentally reimagines the skip connection paradigm of traditional U-Net through its nested and dense skip pathways. Unlike standard U-Net which directly concatenates encoder and decoder features, U-Net++ incorporates intermediate convolutional modules that gradually transform features, effectively bridging the semantic gap and improving segmentation accuracy.

The architectural innovations of U-Net++ address several key limitations of traditional U-Net. First, the semantic gap between encoder and decoder features often leads to suboptimal feature fusion, particularly for small objects like lung nodules. Second, the single skip connection path may not provide sufficient gradient flow during training, potentially limiting the network’s ability to learn fine-grained features essential for accurate nodule boundary delineation.




Fig. 2. U-Net++ architecture with nested skip connections and deep supervision. The nested pathways bridge the semantic gap between encoder and decoder features.

Fig. 2 illustrates the architectural design of U-Net++ used in our system. The network introduces nested skip connections, where encoder features are progressively transformed through intermediate layers before being merged into decoder pathways. This hierarchical transformation alleviates the semantic disparity between low-level and high-level features, facilitating more accurate localization of nodule boundaries.

The network employs deep supervision by producing segmentation outputs at multiple decoder depths. These auxiliary outputs are combined during training to provide more stable gradients, improve learning efficiency, and enhance overall segmentation accuracy. Specifically, the network generates outputs at four different resolution levels: full resolution, 1/2, 1/4, and 1/8 of the input size. The multi-output supervision scheme enables intermediate decoder layers to contribute meaningfully to the final segmentation performance, effectively regularizing the training process and preventing overfitting.

The deep supervision mechanism works by applying the same loss function at each output level and combining them with equal weights. This approach ensures that feature representations at different scales contribute to the learning process, leading to more robust feature extraction and better gradient flow throughout the network. The auxiliary losses help prevent the vanishing gradient problem that can occur in very deep networks.

To address the inherent class imbalance in pulmonary nodule segmentation, where nodule pixels constitute less than 1% of the total image area, U-Net++ employs a hybrid loss function that combines Dice loss and binary cross-entropy:

$$\mathcal{L} = \alpha \mathcal{L}_{Dice} + (1 - \alpha) \mathcal{L}_{BCE} \quad (1)$$

where α is a weighting parameter empirically set to 0.5. The Dice loss is defined as:

$$\mathcal{L}_{Dice} = 1 - \frac{2|P \cap T| + \varepsilon}{|P| + |T| + \varepsilon} \quad (2)$$

where P and T represent predicted and target segmentation masks respectively, and $\varepsilon = 1$ is a smoothing factor to prevent division by zero. The Dice loss captures region overlap and is particularly effective for handling class imbalance, while binary cross-entropy provides pixel-wise classification accuracy. This formulation ensures the model is sensitive to small nodules while maintaining robustness against false positives. The binary cross-entropy (BCE) term is defined as:

$$\mathcal{L}_{BCE} = -\frac{1}{N} \sum_{i=1}^N [t_i \log(p_i + \varepsilon) + (1 - t_i) \log(1 - p_i + \varepsilon)] \quad (3)$$

where N is the number of pixels (voxels), $t_i \in \{0, 1\}$ is the ground-truth label of pixel i , $p_i \in [0, 1]$ is the predicted foreground probability, and ε (set to 10^{-7}) avoids numerical instability. Unlike the Dice term that supervises global region overlap, $\mathcal{L}_{BCE}$ provides fine-grained, voxel-wise supervision and encourages well-calibrated probabilities. Empirically, combining the two losses in Eq. (1) (we use $\alpha = 0.5$) stabilizes optimization under severe class imbalance, yielding sharper boundaries while reducing spurious activations around vessels.

Additionally, we implement a curriculum learning strategy during training, where the model initially focuses on larger, more obvious nodules before gradually incorporating smaller and more challenging cases. This approach helps stabilize training and improves final performance on difficult cases.

2.2 Stage 2: Robust False-Positive Reduction with MGI-CNN

To suppress false positives produced during candidate generation, we employ the Multi-scale Gradual Integration CNN (MGI-CNN), which is specifically designed to handle the complex contextual and morphological variability of pulmonary nodules. This network analyzes volumetric patches of candidate regions at multiple spatial scales and integrates features through a dual-path architecture, enabling it to capture both fine-grained details and broader anatomical context.

The MGI-CNN architecture consists of three main components. The first stage performs multi-scale 3D patch extraction. For each candidate nodule location, the system extracts three volumetric patches at different scales: small ($20 \times 20 \times 6$ voxels), medium ($30 \times 30 \times 10$ voxels), and large ($40 \times 40 \times 26$ voxels). These patches capture complementary information: the smallest scale focuses on local textures and borders, while larger scales include surrounding tissues and spatial context that help distinguish nodules from other anatomical structures. (see Fig. 3)




Fig. 3. Illustration of multi-scale patch extraction and resizing strategy. From the 3D CT volume, patches of varying sizes are extracted around the candidate nodule regions ($40 \times 40 \times 26$, $30 \times 30 \times 10$, $20 \times 20 \times 6$) and resized to a uniform input size of $20 \times 20 \times 6$ to form three resolution levels: S1 (high-resolution), S2 (medium-resolution), and S3 (low-resolution).

In the second stage, these patches are processed through dual pathways: the zoom-in path integrates features from coarse to fine scales, whereas the zoom-out path operates in the reverse direction, combining fine-scale features into broader semantic contexts. This bidirectional fusion enables MGI-CNN to reason about both local nodule morphology and global structure, effectively distinguishing true nodules from confounding anatomical structures such as vessel bifurcations and lymph nodes. (see Fig. 4)




Fig. 4. Illustration of the Gradual Feature Extraction (GFE) mechanism. In the zoom-in stream, features are extracted in a bottom-up manner by concatenating current scale features with those from the previous scale. This hierarchical fusion improves discrimination by leveraging both local and contextual information progressively across S1, S2, and S3.

The zoom-in pathway begins with the largest scale patch and progressively incorporates information from smaller scales. At each integration step, features from the coarser scale are upsampled and concatenated with features from the finer scale, followed by 3D convolutional layers to fuse the multi-scale information. This process allows the network to refine coarse predictions with fine-grained details.

Conversely, the zoom-out pathway starts with the smallest scale and gradually incorporates broader contextual information. This pathway is particularly effective for distinguishing nodules from isolated vessel segments, as it provides the spatial context necessary to understand whether a candidate region represents a true spherical nodule or a cross-sectional view of a cylindrical vessel.

The architectural design of each pathway includes residual connections to facilitate gradient flow and prevent degradation in deeper layers. Batch normalization is applied after each convolutional layer to stabilize training, and dropout with probability 0.3 is used in the fully connected layers to prevent overfitting.

The final stage merges the outputs of both zoom paths through concatenation, applies deep 3D convolutional blocks for feature refinement, and performs binary classification to distinguish true nodules from false positives. The network employs global average pooling to reduce spatial dimensions and prevent overfitting, followed by fully connected layers with dropout regularization. (see Fig. 5)




Fig. 5. MGI-CNN final stage: Convolutional blocks and classification layers for robust nodule discrimination.

2.3 Pipeline Overview and Integration Strategy

The full detection framework is illustrated in Fig. 6. It integrates the strengths of both U-Net++ (for candidate segmentation) and MGI-CNN (for false-positive reduction) in a two-stage cascade. The preprocessing stage involves resampling CT scans to isotropic spacing, intensity normalization, and slice-wise resizing to ensure consistent input dimensions.

First, U-Net++ segments 2D CT slices to generate candidate nodules with high sensitivity. The network processes each slice independently, producing probability maps that highlight potential nodule regions. Connected component analysis is then applied to extract the centroid locations of segmented regions, with morphological operations used to filter out extremely small or elongated components that are unlikely to represent true nodules.

For each candidate centroid, 3D volumes are constructed by extracting patches centered on the 2D coordinates and extending through adjacent slices. These candidate volumes are then fed into the MGI-CNN, which applies multi-scale volumetric analysis to assess the likelihood that each region contains a true nodule.

The network outputs a probability score for each candidate, which is used for final detection decisions.




Fig. 6. Overview of the proposed two-stage lung nodule detection pipeline. The framework seamlessly integrates U-Net++ for high-sensitivity candidate generation with MGI-CNN for robust false-positive reduction.

Post-processing operations include non-maximum suppression to eliminate redundant detections and boundary refinement using the original segmentation masks. This end-to-end cascade ensures a balance between sensitivity and specificity, enabling accurate and reliable detection in complex lung CT scans while maintaining computational efficiency suitable for clinical deployment.

3 Experiments and Results

3.1 Experimental Setup

All experiments were carried out on the LUNA16 dataset, a widely-used benchmark for lung nodule detection that contains 888 CT scans with pixel-level annotations. The dataset provides comprehensive ground truth for nodules larger than 3mm in diameter, making it ideal for evaluating automated detection systems. Subsets 0–2 (267 scans) were used for training and validation, while subset 3 (89 scans) was reserved for testing to ensure unbiased evaluation.

Data preprocessing involved several steps to ensure consistent input quality. Scans were resampled to isotropic 1mm voxels using cubic interpolation to normalize spatial resolution across different scanners. Intensity values were windowed to the range $[-1000, 400]$ HU to focus on lung tissue and nodule appearance, then normalized to $[0, 1]$ for stable training. Each slice was resized to 320×320 pixels to balance computational efficiency with spatial detail preservation.

Ground-truth masks were generated from the official annotations.csv file by drawing circular regions of interest (ROI) with radius $r = \frac{d}{2} + 5$ pixels on intersecting

slices, where d represents the nodule diameter. This approach resulted in 3100 positive masks in the test set, providing sufficient data for reliable performance evaluation.

Two segmentation backbones were implemented and trained: (i) a standard U-Net as baseline and (ii) a U-Net++ with four-way deep supervision as our proposed approach. Both networks were optimized for 80 epochs using the Adam optimizer with learning rate $\eta = 10^{-3}$ and weight decay 10^{-5} . The batch size was set to 16 to balance memory constraints with gradient stability.

A comprehensive data augmentation strategy was employed to improve generalization and robustness. Augmentations included elastic deformation (probability $p = 0.3$), random rotations within $\pm 15^\circ$, horizontal and vertical flips, intensity jitter of $\pm 10\%$, and random translations of ± 15 pixels. Each mini-batch preserved a 1:1 ratio of nodule-positive to nodule-negative slices to mitigate class imbalance effects during training.

3.2 Segmentation Performance Comparison

The training dynamics reveal significant differences between U-Net and U-Net++ architectures. Fig. 7 shows that U-Net quickly overfits after epoch 20, with training accuracy climbing beyond 75% while validation accuracy oscillates around 55%. The validation loss plateaus near 0.55, indicating limited generalization capability.

In contrast, Fig. 8 demonstrates that U-Net++ benefits significantly from nested skip pathways and deep supervision. The validation accuracy stabilizes above 60% and continues to improve gradually, while the validation loss steadily declines below 0.40. This behavior indicates better optimization stability and improved generalization performance.




Fig. 7. Training dynamics for baseline U-Net: **(a)** model accuracy over epochs (train vs. validation), **(b)** model loss over epochs (train vs. validation). The model exhibits clear overfitting behavior after epoch 20, as evidenced by diverging validation curves.




Fig. 8. Training dynamics for U-Net++: **(a)** model accuracy over epochs (train vs. validation), **(b)** model loss over epochs (train vs. validation). The model demonstrates more stable convergence and consistently lower validation loss compared to baseline U-Net.

3.3 Quantitative Segmentation Results

Table 1 summarizes the quantitative segmentation performance on subset 3. While the absolute Dice scores appear modest—30.3% for U-Net and 34.7% for U-Net++—these values are typical for lung nodule segmentation tasks due to the small size of target objects and the challenge of precise boundary delineation. More importantly, the nodule-level recall, which measures the percentage of ground-truth nodules that produce any detectable segmentation response, increases substantially from 93.4% to 97.8%.

Table 1. Quantitative segmentation metrics on LUNA16 subset 3 (slice-level evaluation)

Method	Dice(%)	IoU(%)	HD95(mm)	Recall(%)
Unet	30.3	24.3	25.4	93.4
Unet++	34.7	27.1	21.8	97.8

The improvement in recall is particularly significant because the segmentation stage serves primarily to generate candidates for subsequent classification. Missing nodules during this stage cannot be recovered by the MGI-CNN classifier, making high recall essential for overall system performance. The 4.4 percentage point improvement in recall directly contributes to the downstream classification improvement.

3.4 Visual Segmentation Comparison

Fig. 9 provides qualitative comparison of segmentation results between U-Net and U-Net++ on representative test cases. The visualizations demonstrate that U-Net++ produces more accurate nodule boundaries with fewer fragmented predictions and better adherence to actual nodule shapes.

In challenging cases involving small nodules or those adjacent to blood vessels, U-Net++ demonstrates superior performance in maintaining nodule connectivity while suppressing vessel-related false positives. The nested skip connections enable better preservation of fine-scale features, resulting in more coherent segmentation masks that better reflect true nodule morphology.




Fig. 9. Qualitative comparison of segmentation results. Each row shows the ground truth mask, U-Net prediction, and U-Net++ prediction for a representative case.

3.5 False-Positive Reduction Performance

After connected-component analysis, the segmentation masks produce 75,950 three-dimensional candidate patches (64^3 voxel cubes) from the test subset. Each patch is classified by the MGI-CNN to distinguish true nodules from false positives. Fig. 10 shows the resulting ROC curves for both segmentation backbones.

With proposals generated by the baseline U-Net, the MGI-CNN classifier reaches an AUC of 0.852. Switching to proposals from U-Net++ raises the AUC to 0.901 representing an absolute gain of 4.8 percentage points. This improvement holds consistently across the full operating range, demonstrating that higher-quality candidate masks translate directly into more discriminative 3D features for the classifier.




Fig. 10. Patch-level ROC curves on LUNA16 test subset 3. The improved candidate quality from U-Net++ directly translates to enhanced classification performance.

With proposals generated by the baseline U-Net, the MGI-CNN classifier reaches an AUC of 0.852. Switching to proposals from U-Net++ raises the AUC to 0.901, representing an absolute gain of 4.8 percentage points. This improvement holds consistently across the full operating range, demonstrating that higher-quality candidate masks translate directly into more discriminative 3D features for the classifier.

The performance gain can be attributed to several factors. First, the improved boundary accuracy from U-Net++ provides more representative 3D patches for training and testing the MGI-CNN. Second, the higher recall ensures that fewer true nodules are missed during candidate generation, reducing the bias in the training data. Finally, the more consistent segmentation quality reduces the variability in patch appearance, enabling the classifier to learn more robust feature representations.

3.6 Discussion

The experimental findings yield several important observations that inform both the design and deployment of lung nodule detection systems.

First, we observe that recall during the candidate generation stage has a direct and significant impact on overall detection performance. Even a modest 4.4% increase in nodule-level recall (from 93.4% to 97.8%) leads to a 4.9-point gain in classification AUC (from 0.852 to 0.901). This highlights that maximizing sensitivity in the segmentation stage is more critical than achieving pixel-perfect boundaries, as missed nodules cannot be recovered downstream.

Second, the integration of deep supervision in U-Net++ demonstrably improves training stability and generalization. Compared to the baseline U-Net, U-Net++ achieves lower validation loss and approximately 6% higher validation accuracy. This suggests that nested skip connections and intermediate supervision signals help regularize training and mitigate overfitting—an especially valuable property when working with limited annotated medical datasets.

Third, the architectural benefits of U-Net++, originally proposed for general medical imaging tasks, transfer effectively to the more domain-specific challenge of

pulmonary nodule segmentation. This supports the broader applicability of its design innovations, including semantic feature bridging and dense pathway learning.

Limitations and Future Work: While the proposed approach yields strong results, certain limitations persist. The segmentation module operates on 2D slices independently, potentially overlooking inter-slice continuity crucial for detecting subtle nodules. A natural extension is to explore fully 3D segmentation architectures that better exploit volumetric coherence. Incorporating attention mechanisms may further refine feature learning, particularly for detecting sub-6mm nodules with ground-glass opacity—one of the most challenging cases.

From a system-level perspective, future research could investigate end-to-end optimization of the segmentation and classification stages to reduce error propagation. Additionally, integrating vision transformers or leveraging self-supervised pretraining on large-scale unlabeled scans may offer further improvements. Clinical validation across diverse patient cohorts, scanner protocols, and institutions will be essential to ensure the robustness and generalizability of the proposed framework.

4 Conclusion

4.1 Key Contributions and Achievements

Our work makes several important contributions to the field of automated medical image analysis.

First, we introduce an enhanced two-stage detection pipeline that leverages the architectural innovations of U-Net++ for superior candidate generation. The nested skip pathways and deep supervision mechanism result in a substantial improvement in nodule-level recall from 93.4% to 97.8%, directly addressing the critical challenge of missing small nodules during the initial detection phase.

Second, we provide comprehensive empirical evidence that improved segmentation quality in the first stage directly translates to enhanced classification performance in the second stage. The integration of U-Net++ with MGI-CNN achieves a classification AUC improvement from 0.852 to 0.901, representing a significant advancement in false-positive reduction capabilities.

Third, our detailed analysis of training dynamics reveals that U-Net++’s architectural design provides superior optimization stability and generalization performance compared to traditional U-Net approaches. This finding has broader implications for medical image segmentation tasks where training data may be limited and model robustness is paramount.

4.2 Clinical Significance and Impact

The proposed framework addresses several critical challenges in clinical lung cancer screening programs. The improved recall rate of 97.8% means that fewer than 3% of nodules are missed during the initial detection phase, substantially reducing the risk of overlooking potentially malignant lesions. Simultaneously, the enhanced specificity (AUC = 0.901) significantly reduces false-positive rates, thereby minimizing unnecessary patient anxiety and reducing healthcare costs associated with follow-up procedures.

4.3 Methodological Insights

Our research provides valuable insights into the design principles for effective medical image analysis pipelines. The success of the U-Net++ and MGI-CNN integration demonstrates the importance of architectural synergy in multi-stage detection systems. Specifically, we show that the quality of candidate generation in the first stage has cascading effects on the overall system performance, emphasizing the need for high-quality segmentation networks in two-stage pipelines.

The bidirectional feature integration approach employed by MGI-CNN proves particularly effective for distinguishing true nodules from anatomical confounders. The zoom-in and zoom-out pathways capture both fine-grained textural details and broader spatial context, enabling robust discrimination between nodules and vessel cross-sections—a common source of false positives in automated detection systems.

4.4 Future Research Directions

Several promising avenues for future research emerge from our work. **Immediate extensions** include the development of fully 3D segmentation approaches that can better leverage inter-slice continuity information, potentially further improving detection accuracy for small nodules. The integration of attention mechanisms could provide more targeted feature learning, particularly beneficial for detecting subtle ground-glass opacities.

Long-term research directions encompass the exploration of end-to-end optimization strategies that jointly train both segmentation and classification networks, potentially reducing error propagation between pipeline stages. The incorporation of emerging technologies such as vision transformers and self-supervised learning represents another promising direction for continued advancement in automated lung nodule detection.

Clinical validation remains a critical next step, requiring extensive testing across diverse patient populations, scanner types, and imaging protocols to establish robustness and generalizability. Multi-institutional studies would provide valuable insights into the system’s performance across different clinical environments and patient demographics.

4.5 Concluding Remarks

The integration of advanced deep learning architectures in medical image analysis continues to show tremendous promise for improving clinical outcomes. Our work demonstrates that thoughtful architectural design and systematic integration of complementary technologies can yield substantial performance improvements while maintaining practical viability. The synergistic combination of U-Net++ and MGI-CNN represents a significant step forward in automated lung nodule detection, providing a robust foundation for future developments in computer-aided diagnosis systems.

As the field moves toward more sophisticated AI-assisted medical imaging tools, the principles demonstrated in this work—architectural innovation, systematic integration, and comprehensive validation—will remain fundamental to developing effective and clinically viable solutions. The continued advancement of such technologies holds great promise for enhancing early cancer detection capabilities and ultimately improving patient outcomes worldwide.

References

1. Bray, F., Laversanne, M., Sung, H., et al.: Global cancer statistics 2022: GLOBOCAN estimates of incidence and mortality worldwide for 36 cancers in 185 countries. *CA: A Cancer Journal for Clinicians* 74(3), 229–263 (2024).
2. Ronneberger, O., Fischer, P., Brox, T.: U-Net: Convolutional Networks for Biomedical Image Segmentation. In: Navab, N., Hornegger, J., Wells, W.M., Frangi, A.F. (eds.) *MICCAI 2015. LNCS*, vol. 9351, pp. 234–241. Springer, Cham (2015).
3. Zhou, Z., Siddiquee, M.M.R., Tajbakhsh, N., Liang, J.: UNet++: A Nested U-Net Architecture for Medical Image Segmentation. In: Frangi, A.F., Schnabel, J.A., Davatzikos, C., Alberola-López, C., Fichtinger, G. (eds.) *Deep Learning in Medical Image Analysis and Multimodal Learning for Clinical Decision Support. LNCS*, vol. 11045, pp. 3–11. Springer, Cham (2018).
4. Kim, B.-C., Choi, J.-S., Suk, H.-I.: Multi-scale gradual integration CNN for false positive reduction in pulmonary nodule detection. *Neural Networks* 115, 1–10 (2019).
5. Gao, C., Wu, L., Wu, W., Huang, Y., Wang, X., Sun, Z., Xu, M., Gao, C.: Deep learning in pulmonary nodule detection and segmentation: A systematic review. *European Radiology* 35, 255–266 (2025).
6. Li, Y., Yan, B., He, S.: Advances and challenges in the treatment of lung cancer. *Biomedicine & Pharmacotherapy* 169, Article 115891 (2023).
7. Adams, S.J., Stone, E., Baldwin, D.R., Vliegenthart, R., Lee, P., Fintelmann, F.J.: Lung cancer screening. *The Lancet* 401, 390–408 (2023).
8. Zhao, D., Liu, Y., Yin, H., Wang, Z.: An attentive and adaptive 3D CNN for automatic pulmonary nodule detection in CT image. *Expert Systems with Applications* 211, Article 118672 (2023).
9. Huang, Y.-S., Chou, P.-R., Chen, H.-M., Chang, Y.-C., Chang, R.-F.: One-stage pulmonary nodule detection using 3-D DCNN with feature fusion and attention mechanism in CT image. *Computer Methods and Programs in Biomedicine* 220, Article 106786 (2022).
10. Wu, Y., Zhou, Y., Zeng, W., Qian, Q., Song, M.: An attention-based 3D CNN with multi-scale integration block for Alzheimer’s disease classification. *IEEE Journal of Biomedical and Health Informatics* 26(11), 5665–5673 (2022).
11. Dinh, T.L., Lee, S.-H., Kwon, S.-G., Kwon, K.-R.: Cell nuclei segmentation in CryoNuSeg dataset using nested U-Net with EfficientNet encoder. In: *Proc. IEEE Int. Conf. on Electronics, Information, and Communication (ICEIC)*, pp. 1–4. IEEE, Jeju, Korea (2022).
12. Huang, L., Miron, A., Hone, K., Li, Y.: Segmenting medical images: From UNet to Res-UNet and nnUNet. In: *Proc. IEEE 37th Int. Symposium on Computer-Based Medical Systems (CBMS)*, pp. 483–489. IEEE, Guadalajara, Mexico (2024).
13. Rayed, M.E., Islam, S.M.S., Niha, S.I., Jima, J.R., Kabir, M.M., Mridha, M.F.: Deep learning for medical image segmentation: State-of-the-art advancements and challenges. *Informatics in Medicine Unlocked* 47, Article 101504 (2024).
14. Azad, R., Aghdam, E.K., Rauland, A., Jia, Y., Avval, A.H., Bozorgpour, A., Karimijafarbigloo, S., Cohen, J.P., Adeli, E., Merhof, D.: Medical image segmentation review: The success of U-Net. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 46(12), 10076–10095 (2024).
15. Bhattacharjee, A., Murugan, R., Goel, T., Mirjalili, S.: Pulmonary nodule segmentation framework based on fine-tuned and pretrained deep neural network using CT images. *IEEE Transactions on Radiation and Plasma Medical Sciences* 7(4), 394–404 (2023).

16. Zhou, Z., Gou, F., Tan, Y., Wu, J.: A cascaded multi-stage framework for automatic detection and segmentation of pulmonary nodules in developing countries. *IEEE Journal of Biomedical and Health Informatics* 26(11), 5619–5630 (2022).